



# Monitoring with Prometheus

# Agenda

---

Observability meaning

---

Difference between (logs,traces,metrics)

---

Metrics types

---

What is Prometheus

---

Prometheus architecture

---

Prometheus installing ( prometheus, node exporter, client library)

---

Configuring prometheus.yaml

---

PromQL




# What is Observability

---

# What is Observability

- **Observability**- the ability to understand and measure the state of a system based upon data generated by the system
- Observability allows you to generate actionable outputs from **unexpected** scenarios
- Observability will help:
  - Give better insight into the internal working of a system/application
  - Speed up troubleshooting
  - Detect hard to catch problems
  - Monitor performance of an application
  - Improve cross-team collaboration


The main purpose of observability is to better understand the internals of your system



# How do we accomplish observability ?



logging



Metrics



traces

```
james@ilmiontdesktop:~$ k logs heron-web-57d58889b-bpvlv --tail=9
```

```
10.244.0.76 - - [20/Oct/2021:18:29:33 +0000] "GET /robots.txt HTTP/1.1" 200 3
10.244.0.76 - - [20/Oct/2021:18:29:34 +0000] "GET /assets/james.jpg HTTP/1.1"
10.244.0.76 - - [20/Oct/2021:18:29:35 +0000] "GET /people/james HTTP/1.1" 304
10.244.0.76 - - [20/Oct/2021:18:34:00 +0000] "GET /robots.txt HTTP/1.1" 200 3
10.244.0.76 - - [20/Oct/2021:18:34:02 +0000] "GET /privacy HTTP/1.1" 200 3427
10.244.0.76 - - [20/Oct/2021:18:49:48 +0000] "GET /robots.txt HTTP/1.1" 200 3
10.244.0.76 - - [20/Oct/2021:19:03:31 +0000] "GET /wordpress/ HTTP/1.1" 301 2
10.244.0.76 - - [20/Oct/2021:19:03:31 +0000] "GET /wordpress HTTP/1.1" 404 24
10.244.0.76 - - [20/Oct/2021:19:31:53 +0000] "GET /robots.txt HTTP/1.1" 200 3
```

## logging

- Logs are records of events that have occurred and encapsulate information about the specific event
- Logs are comprised of :
  - Timestamp of when the log occurred
  - Message containing information

# Traces

---

Traces allow you to follow operations as they traverse through various systems & services

---

So we can follow an individual request and see it flow through our system hop by hop

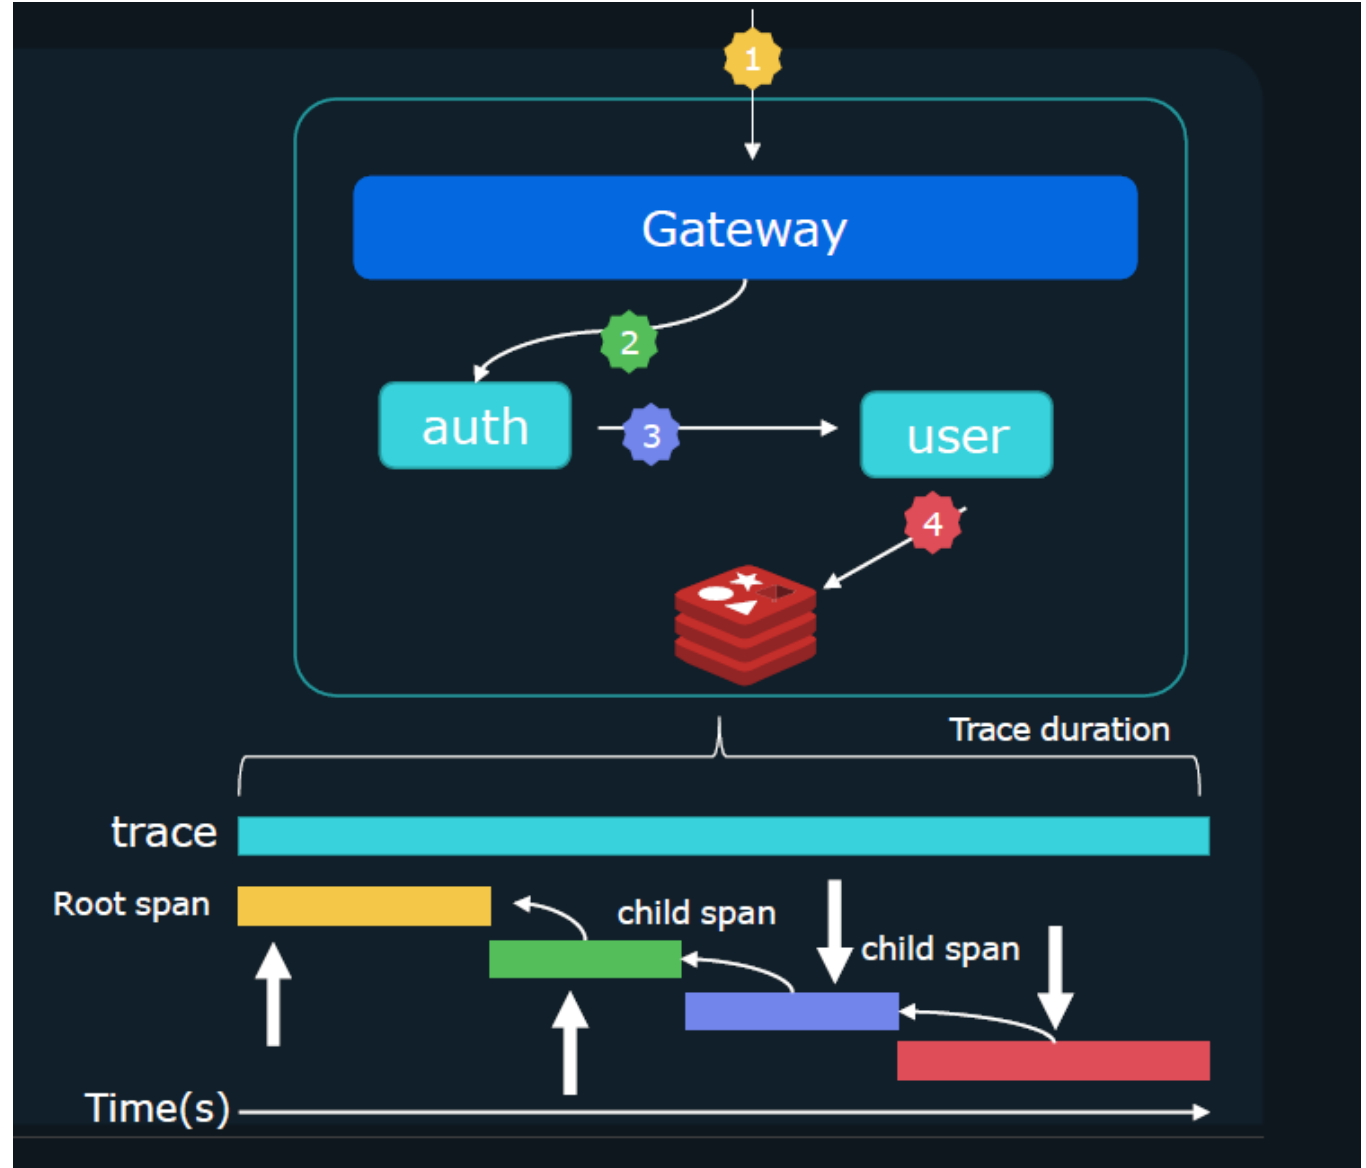
---

Traces help us connect the dots on how processes and services work together



# traces

- Each trace has a trace-id that can be used to identify a request as it traverses the system
- Individual events forming a trace are called spans
- Each span tracks the following:
  - Start time
  - Duration
  - Parent -id



# Metrics

---

- Metrics provide information about the state of a system using numerical values
  - CPU Load
  - Number of open files
  - HTTP response times
  - Number of errors

The data collected can be aggregated over time and graphed using visualization tools to identify trends over time

# Prometheus



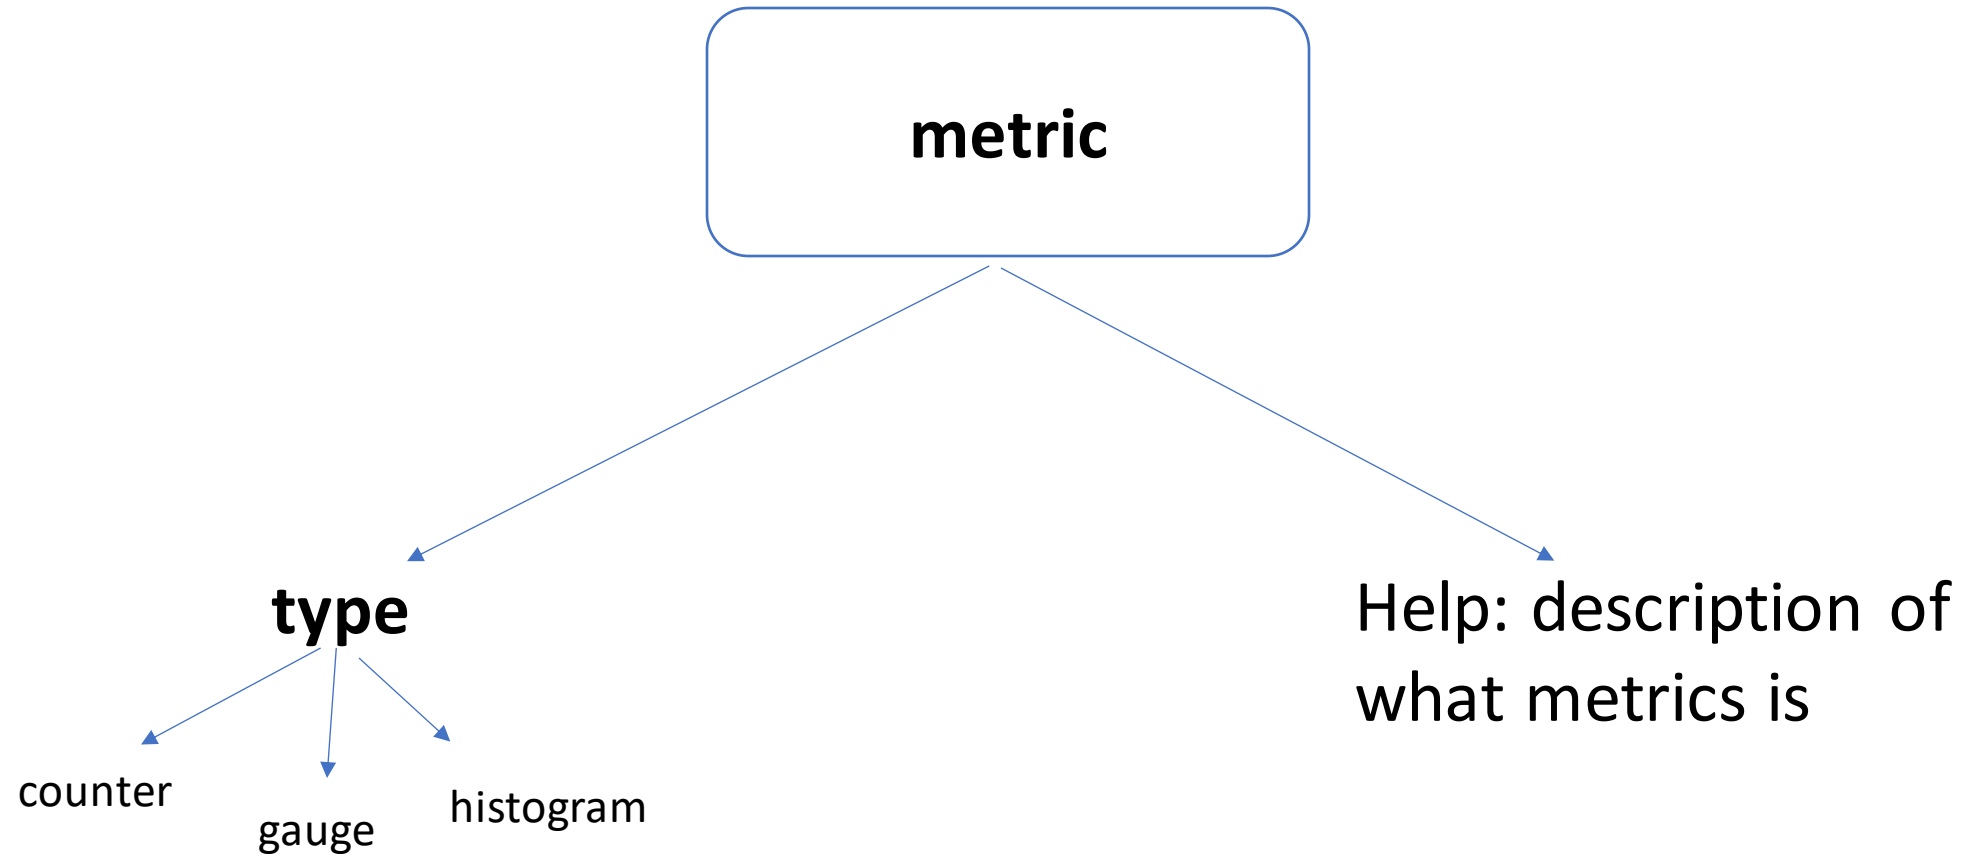
logs

Metrics

Traces

Prometheus is a monitoring solution that is responsible for collecting and aggregating **Metrics**

# Metric types



# Metric type: Counter



```
http_requests_total 10000000
```

```
cpu_seconds_total 3000
```

@ 22:00

```
http_requests_total 970000
```

```
cpu_seconds_total 3000
```

@ 21:00

```
http_requests_total - http_requests_total offset 1h  
= 300000
```

# Metric type: Gauge



`http_requests_active` **2000**

`memory_allocated_bytes` **4.832e+09**

@ 22:00

`http_requests_active` **900**

`memory_allocated_bytes` **3.642e+09**

@ 21:00



`memory_allocated_bytes / (1024*1024*1024)`  
`= 4.5      # gigabytes`

# Metric type: Histogram



```
calculation_seconds_bucket{le="1"}    0  
calculation_seconds_bucket{le="5"}    3  
calculation_seconds_bucket{le="10"}   6  
calculation_seconds_bucket{le="20"}   9  
calculation_seconds_bucket{le="60"}  10
```

@ 21:00



```
calculation_seconds_bucket{le="20"} /  
calculation_seconds_bucket{le="+Inf"} # SLA
```

# What is Prometheus

---

- [Prometheus](#) is an open-source systems monitoring and alerting toolkit originally built at [SoundCloud](#).
- Since its inception in 2012, many companies and organizations have adopted Prometheus, and the project has a very active developer and user [community](#).
- It is now a standalone open source project and maintained independently of any company. To emphasize this, and to clarify the project's governance structure, Prometheus joined the [Cloud Native Computing Foundation](#) in 2016 as the second hosted project, after [Kubernetes](#).
- Prometheus collects and stores its metrics as time series data, i.e. metrics information is stored with the timestamp at which it was recorded, alongside optional key-value pairs called labels.

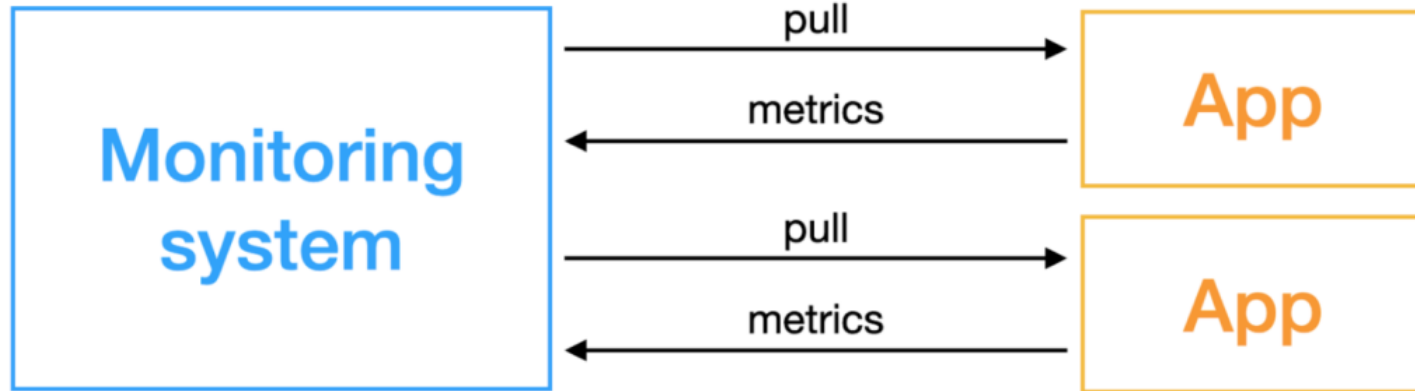


# Why using Prometheus


---

- **Multidimensional data model** – Using time-series data, which is identified by metric name and key-value pairs.
- **PromQL** – A flexible querying language that can leverage the multi-dimensional data model.
- **No reliance on distributed storage** – All single server nodes remain autonomous.
- **Pull model** – Prometheus can collect time-series data by actively “pulling” data over HTTP.
- **Pushing time-series data** – Available through the use of an intermediary gateway.
- **Monitoring target discovery** – Available through static configuration or service discovery.
- **Visualization** – Prometheus offers multiple types of graphs and dashboards.


## Pull-based system



## Push-based monitoring system

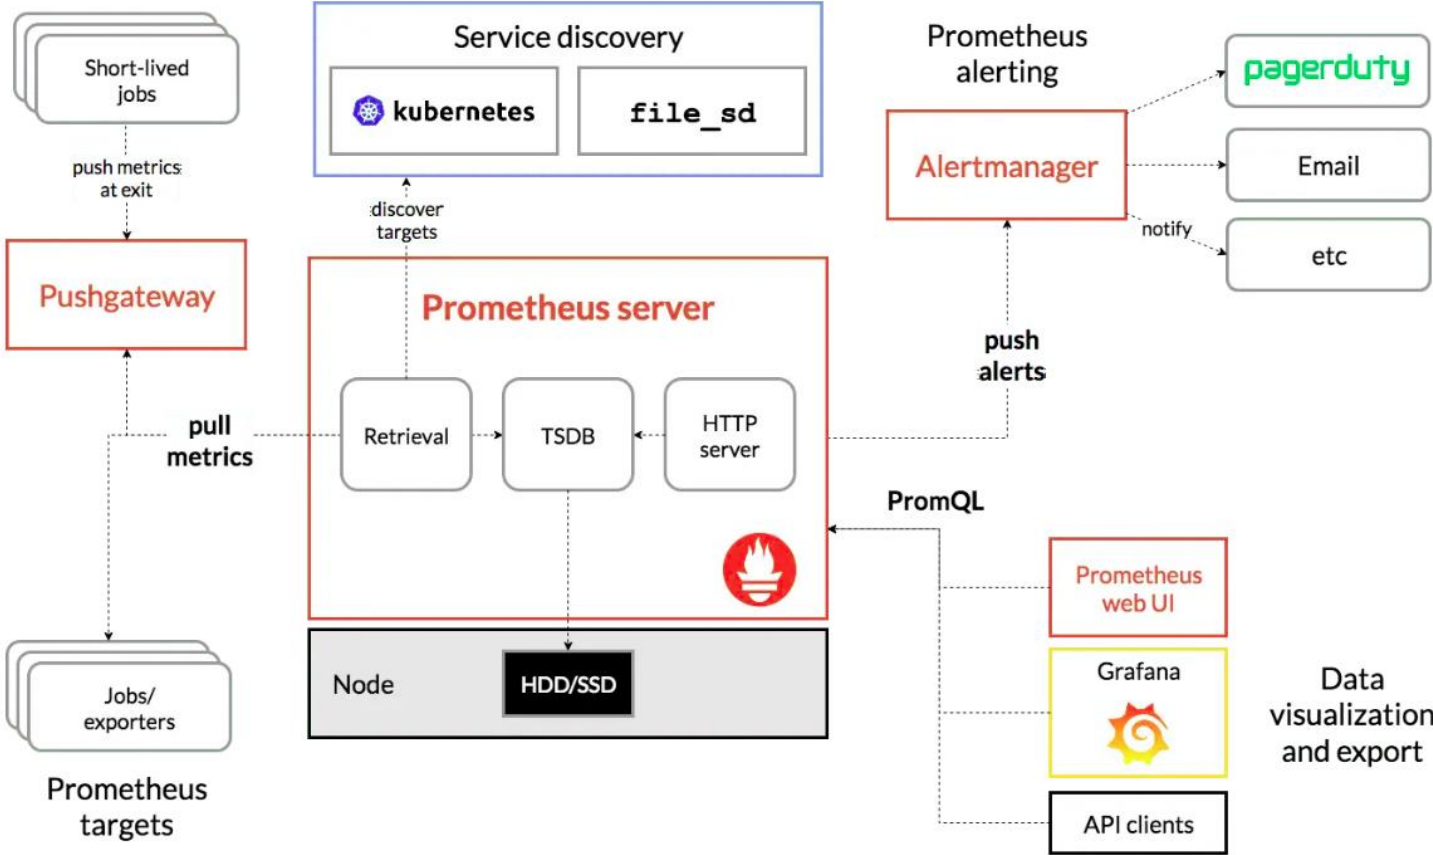


# Prometheus Architecture




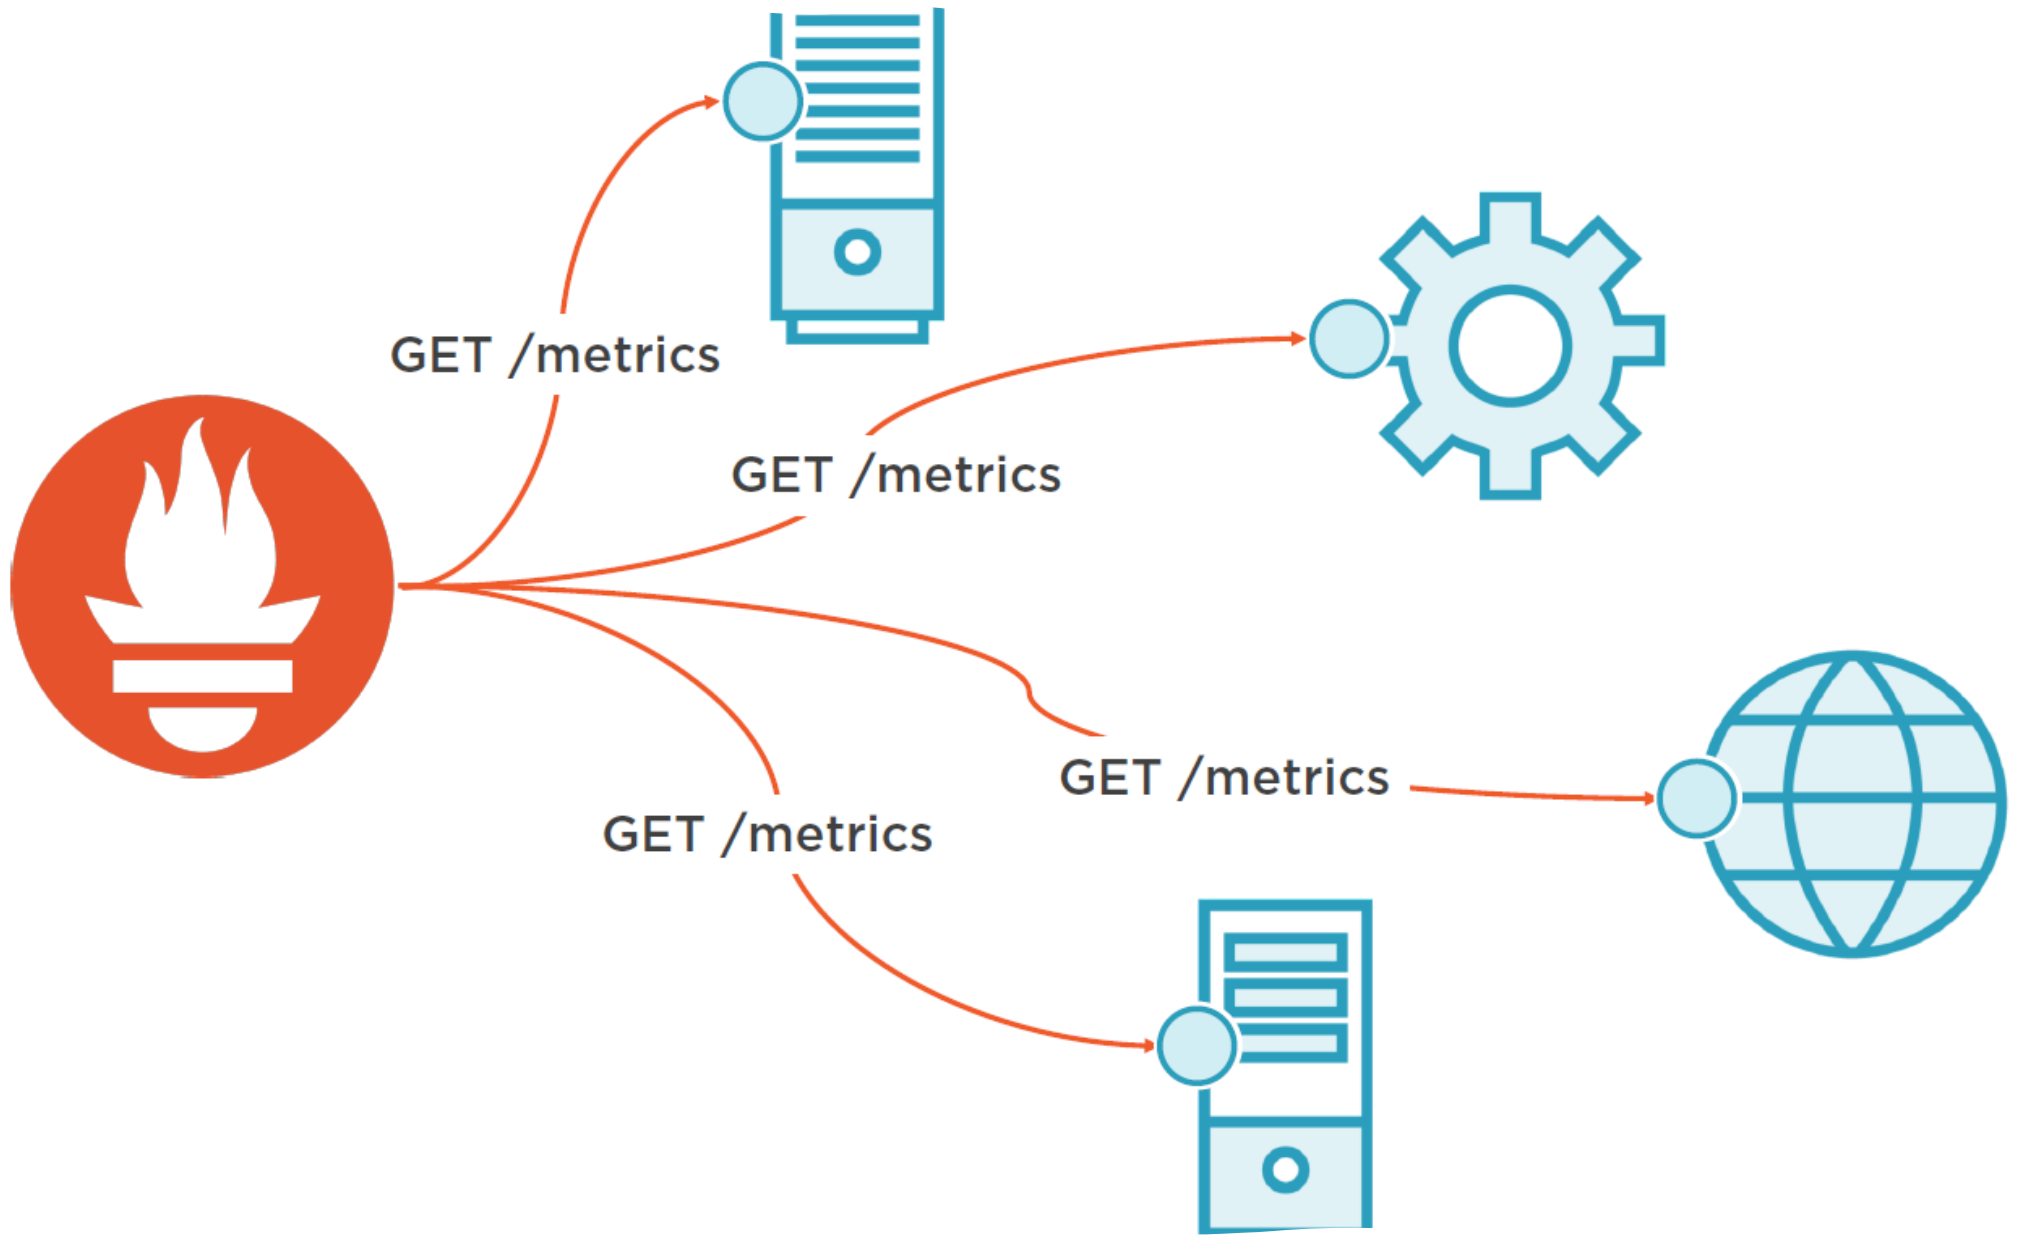
# Prometheus Architecture


Here is the high-level architecture of Prometheus.



source: prometheus.io








```
scrape_configs:
```

- job\_name: 'linux'  
static\_configs:
  - targets: ['ub1804:9100']
- job\_name: 'batch'  
static\_configs:
  - targets: ['ub1804:8080']
- job\_name: 'windows'  
static\_configs:
  - targets: ['win2019:9182']
- job\_name: 'web'  
static\_configs:
  - targets: ['win2019:8080']





# Demo



## Running Prometheus

- Download options
- Running the server
- Exploring the UI

# Labels




```
http_requests_total{code="200",path="/" } 800
http_requests_total{code="500",path="/p1" } 12980
http_requests_total{code="500",path="/p2" } 1064
http_requests_total{code="404",path="/p3" } 36
```



```
sum without(code, path) (http_requests_total)
= 14880      # all requests
```


```
sum without(path) (http_requests_total{code="500"})
= 14044      # all errors
```



### Client libraries



### Exporters



## file-sd.yml

global:

scrape\_interval: 15s

scrape\_configs:

- job\_name: prometheus  
static\_configs:
  - targets: [localhost:9090]
- job\_name: web-file  
file\_sd\_configs:
  - files:
    - 'web.json'

## web.json

```
[  
    {  
        "targets": [ "win2019:9182" ]  
    },  
    {  
        "targets": [ "win2019:8080" ]  
    }  
]
```



# Exploring PromQL Syntax

---



http\_request\_duration\_seconds

```
http_request_duration_seconds
{ job="api", instance="01" }
http_request_duration_seconds
{ job="api", instance="02" }
http_request_duration_seconds
{ job="web", instance="01" }
http_request_duration_seconds
{ job="web", instance="02" }
```



http\_request\_duration\_seconds



```
worker_jobs_total  
{instance="i1",  
  status="processed"} 150
```

```
sum  
without(instance, status)  
(worker_jobs_total)
```



| instance | job_status | job_count |
|----------|------------|-----------|
| i1       | processed  | 150       |

```
SELECT  
SUM(job_count) FROM  
job_summaries
```



170K rows /24hr

- per status
- per instance

~ 50b per row

| instance | job_status | job_count | timestamp  |
|----------|------------|-----------|------------|
| i1       | processed  | 150       | 1592210327 |





```
worker_jobs_total  
{instance="i1",  
  status="processed"}
```

```
150 @ 1592210327  
158 @ 1592210357  
210 @ 1592210387  
235 @ 1592210417
```

One time series < 2b per sample

- per status
- per instance



```
histogram_quantile(  
  0.90,  
  sum without(code,instance)(  
    rate(http_request_seconds[5m]  
  )))
```



```
select top (1) percentile_cont(0.90)  
  within group (order by avg_duration)  
    over () as percentile_90  
from (select avg(duration) as avg_duration,  
            percentile_cont(0.90)  
              over (order by avg(duration))  
              as percentile_90  
from t  
  group by status_code, instance  
) t;
```

worker\_jobs\_active



```
worker_jobs_active
  {instance="i1", job="batch"} 84
worker_jobs_active
  {instance="i2", job="batch"} 51
```

worker\_jobs\_active  
**{instance="i1"}**



```
worker_jobs_active
  {instance="i1", job="batch"} 84
```

worker\_jobs\_active  
**{job="batch", instance=~"i.\*"}**



```
worker_jobs_active
  {instance="i1", job="batch"} 84
worker_jobs_active
  {instance="i2", job="batch"} 51
```

`worker_jobs_active[3m]`



`worker_jobs_active`

`{instance="i1", job="batch"}`

**70** @1592319615.353

**19** @1592319675.357

**34** @1592319735.352

`worker_jobs_active`

`{instance="i2", job="batch"}`

**95** @1592319645.816

**56** @1592319705.818

**55** @1592319765.823

worker\_jobs\_active

{instance="i1"}

[3m]



worker\_jobs\_active

{instance="i1", job="batch"}

**70** @1592319615.353

**19** @1592319675.357

**34** @1592319735.352

```
sum(worker_jobs_active)
```

135

```
sum without(job)  
(worker_jobs_active)
```

worker\_jobs\_active  
 {instance="i1"} 80

worker\_jobs\_active  
 {instance="i2"} 55

**delta(worker\_jobs\_active[1h])**



worker\_jobs\_active

{instance="i1"} **18.305058891159**

worker\_jobs\_active

{instance="i2"} **8.1355748348591**

**avg(delta(worker\_jobs\_active[1h]))**



**13.220316863009444**

```
rate(worker_jobs_total[5m])
```



```
worker_jobs_active
```

```
{instance="i1",status="p"} 47.4
```

```
worker_jobs_active
```

```
{instance="i1 ",status="f"} 4.9
```

```
worker_jobs_active
```

```
{instance="i2",status="p"} 46.2
```

```
worker_jobs_active
```

```
{instance="i2 ",status="f"} 4.7
```

```
sum(rate(worker_jobs_total[5m]))
```



```
103.43267213350337
```



Operator

Metric name


Range

```
sum(rate(worker_jobs_total{instance="i1"}[5m]))
```

Function

Selector

sum without(job, os, runtime) (worker\_jobs\_active)



# Summary



- 
- Prometheus architecture
  - Monitoring systems
  - Metric format
  - Running Prometheus
  - Configuring Prometheus
  - Target labels
  - PromQL
  - Typical expressions

# Q&A



# Thank you

---

